High Performance Computing for Weather and Climate (HPC4WC)

Content: Distributed Memory Parallelism / MPI
Lecturers: Oliver Fuhrer
Block course 701-1270-00L
Summer 2025

Learning Goals

- Understand distributed memory parallelism and how it is different from shared memory parallelism
- Learn basic message passing patterns using MPI
- Be able to apply domain decomposition for solving partial differential equations
- Understand the concept of halo points and able to implement a halo-update.

Computer Architecture

Von Neumann

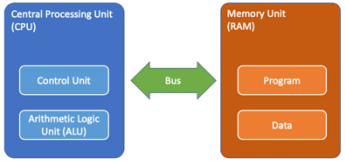

Cache hierarchy

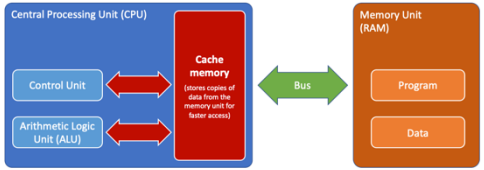

Modern CPU

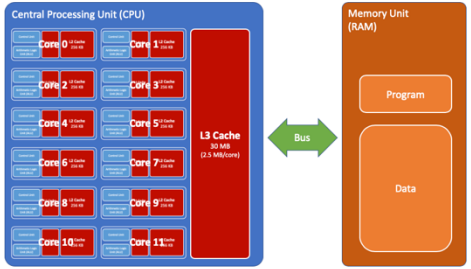

Multicore Node

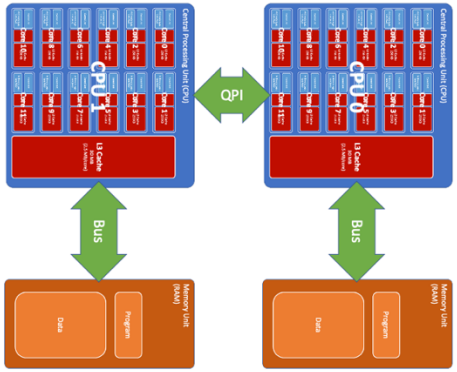

Many nodes

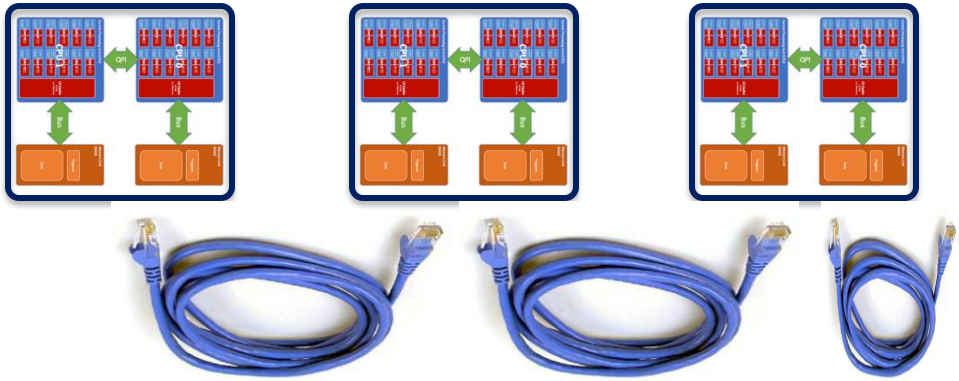

Shared vs. Distributed Memory

All cores on a node share the same address space / memory

```
>>> print(a)
42
```

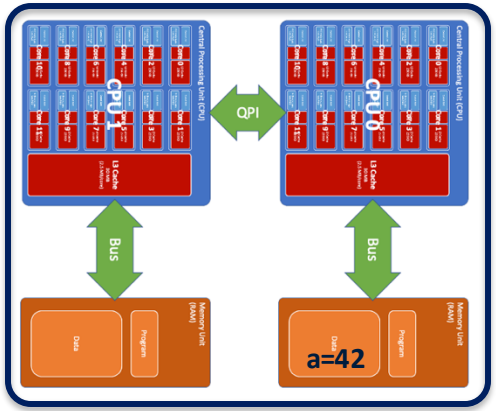

Nodes have different address spaces / memories. Variables are not shared.

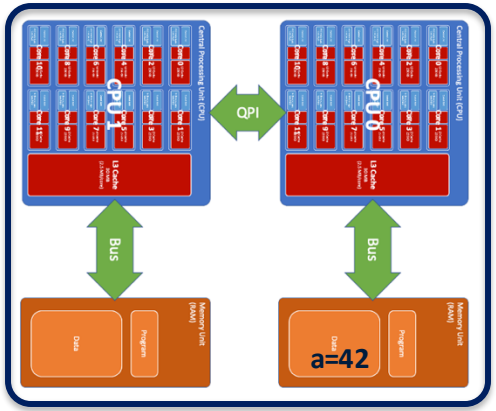


```
>>> print(a)
NameError: name 'a' is not defined
```

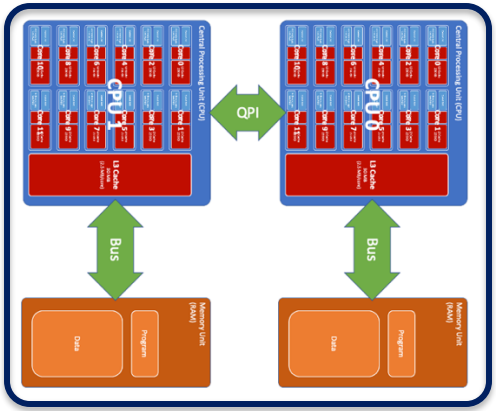

Supercomputer Architecture

(Numbers are for ALPS GH200 and vary from system to system)

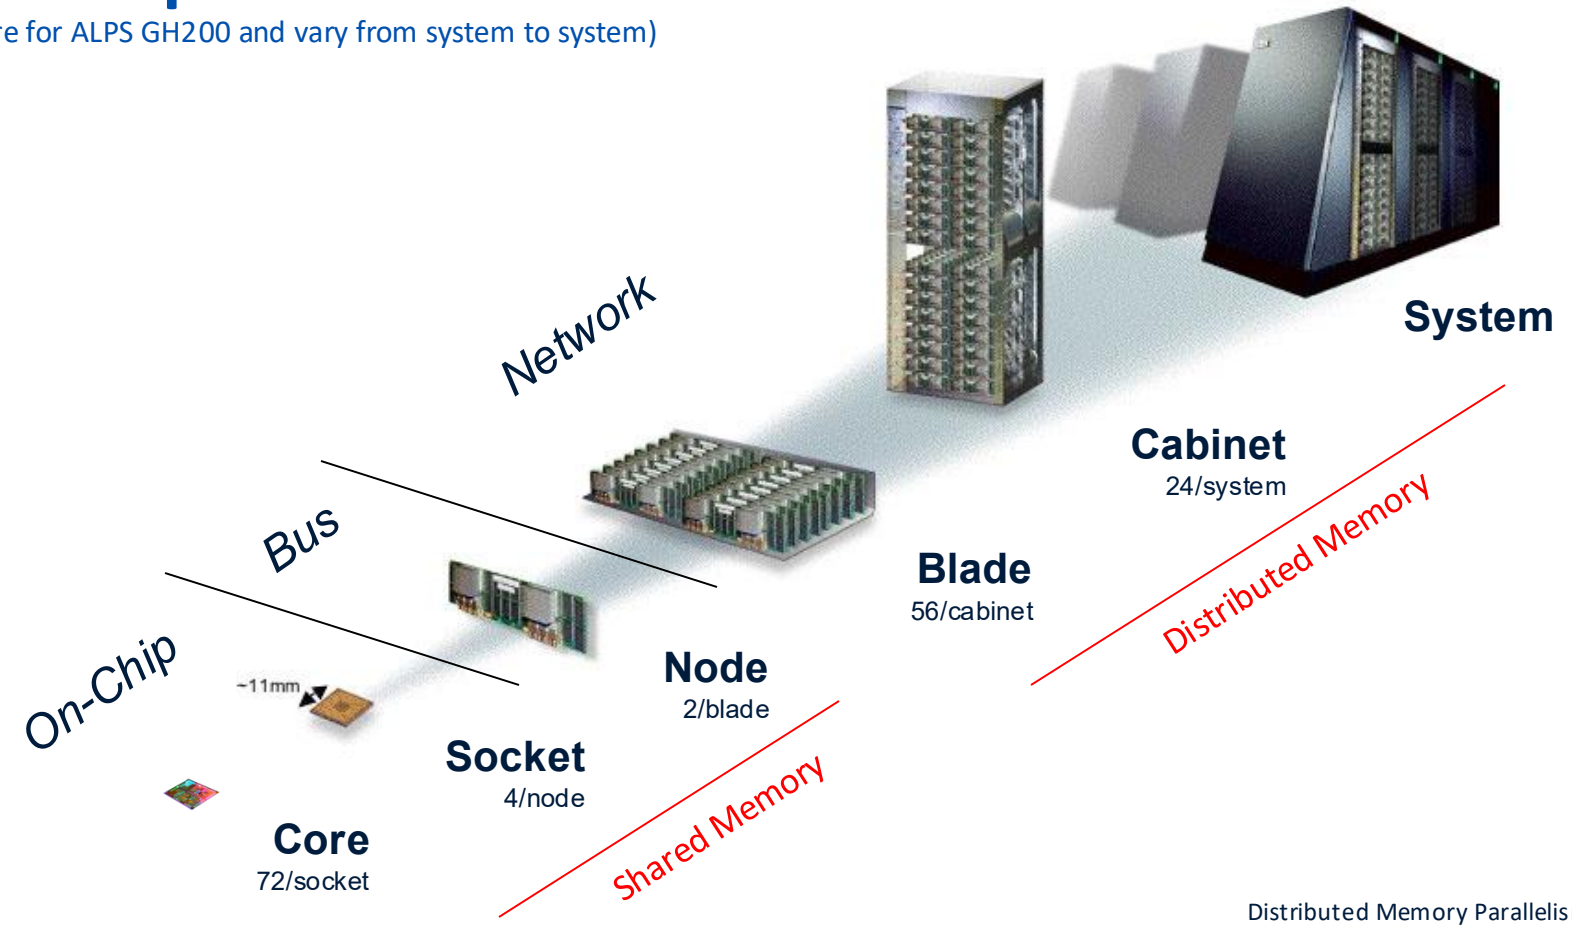

Message Passing

- Information between nodes is transferred over a network cable using a message passing protocol.

```
>>> a = 42
>>> address(a)
0x001a947e3211
>>> send(a, destination=1)
```

```
>>> a = recv(source=0)
>>> print(a)
42
>>> address(a)
0x002f33498e77
```

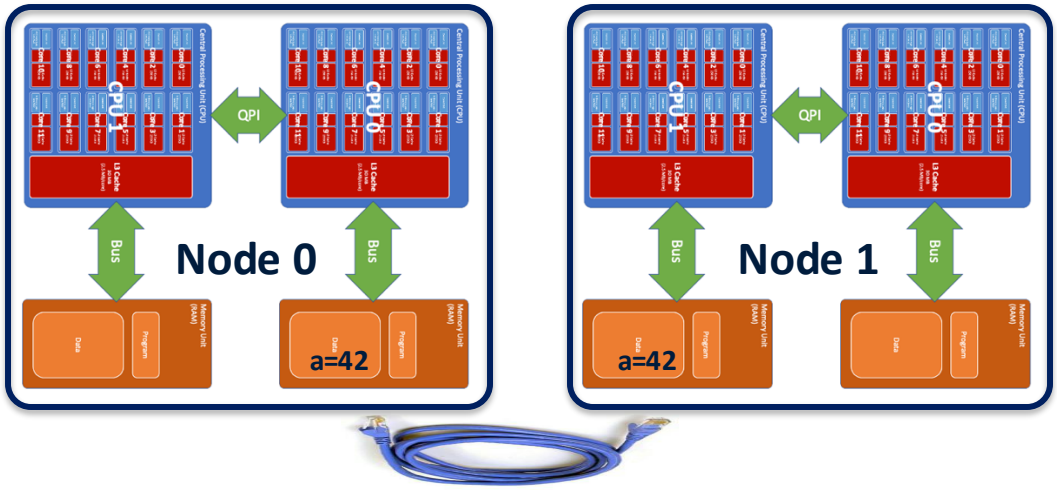

Domain Decomposition

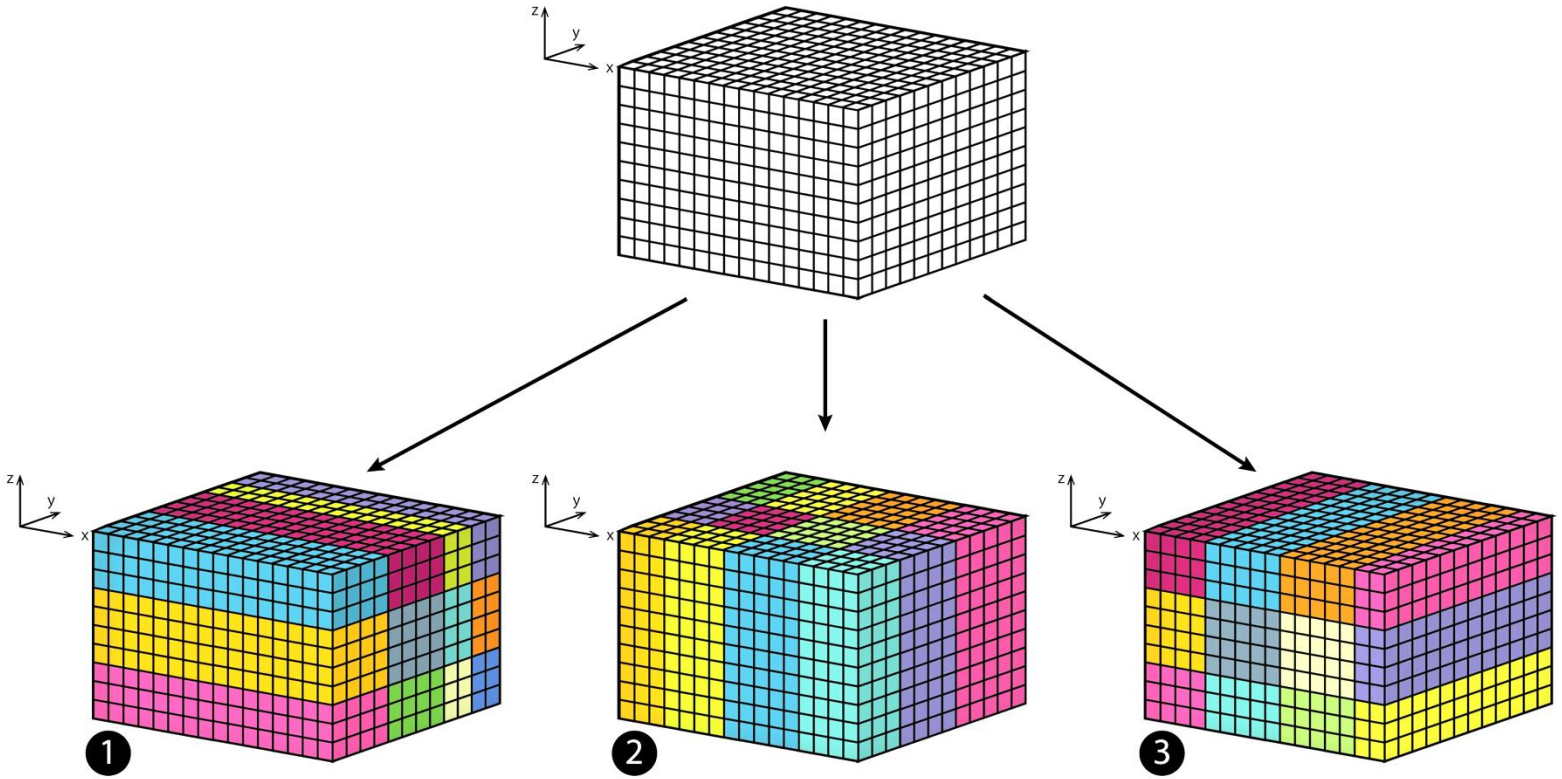

Domain Decomposition in Atmospheric Models

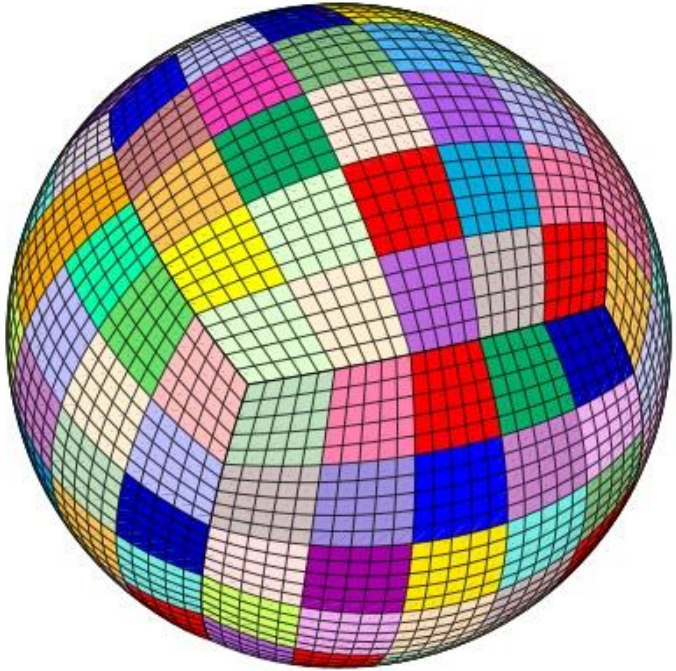

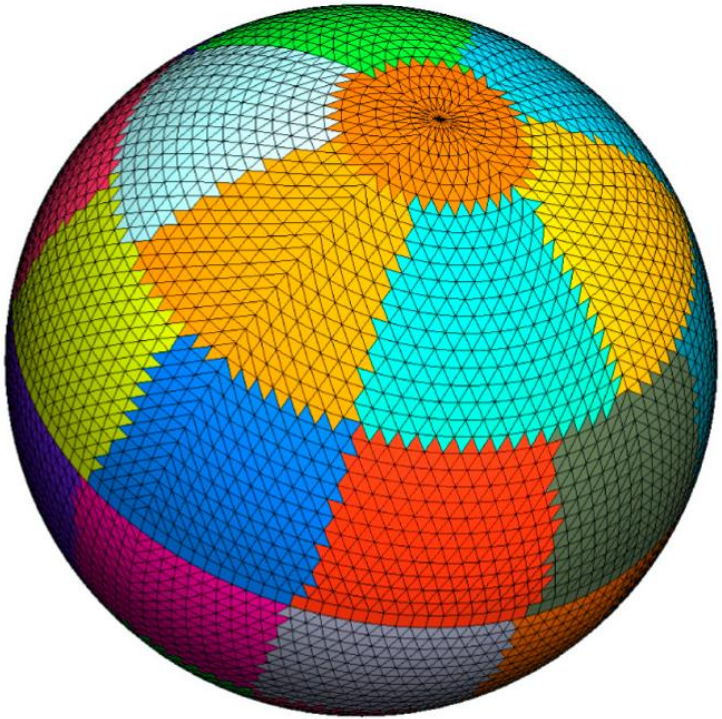

Message Passing Interface (MPI)

- MPI is a standardized and portable message passing standard. (https://www.mpi-forum.org/ and https://github.com/mpi-forum)
- Version 1.0 in 1992, latest Version 5.0 in June 2025
- Support for Fortran, C, C++, Python, Julia, …
- Implemented as a library that provides message passing semantics.
- Several implementations
  - MVAPICH
  - OpenMPI
  - Cray MPI
  - …
- Available on almost any architecture
  - Linux Laptop (apt-get install mpich)
  - Supercomputer
  - Google Cloud Platform
  - …

Lab Exercises

01-test-MPI-setup.ipynb

- Test the setup of your JupyterHub Server to make sure that MPI is working correctly.

02-MPI-introduction.ipynb

- Step-by-step introduction to MPI concepts in Python (mpi4py).

03-domain-decomposition.ipynb

- Learn about domain-decomposition.
- Apply domain-decomposition to a simple 1d example.
- Apply domain-decomposition to the stencil2d.py program.

Send / Receive (Rendezvous protocol = large messages)

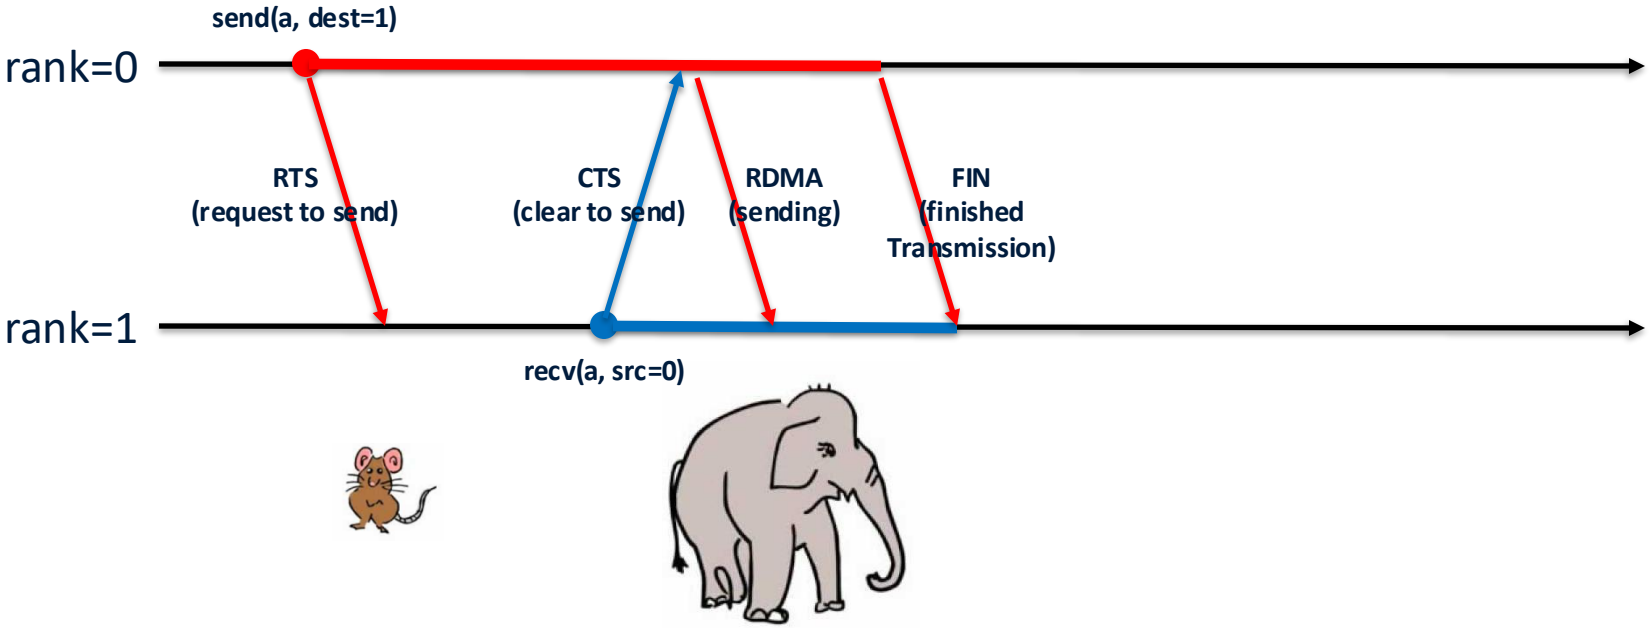

Deadlock

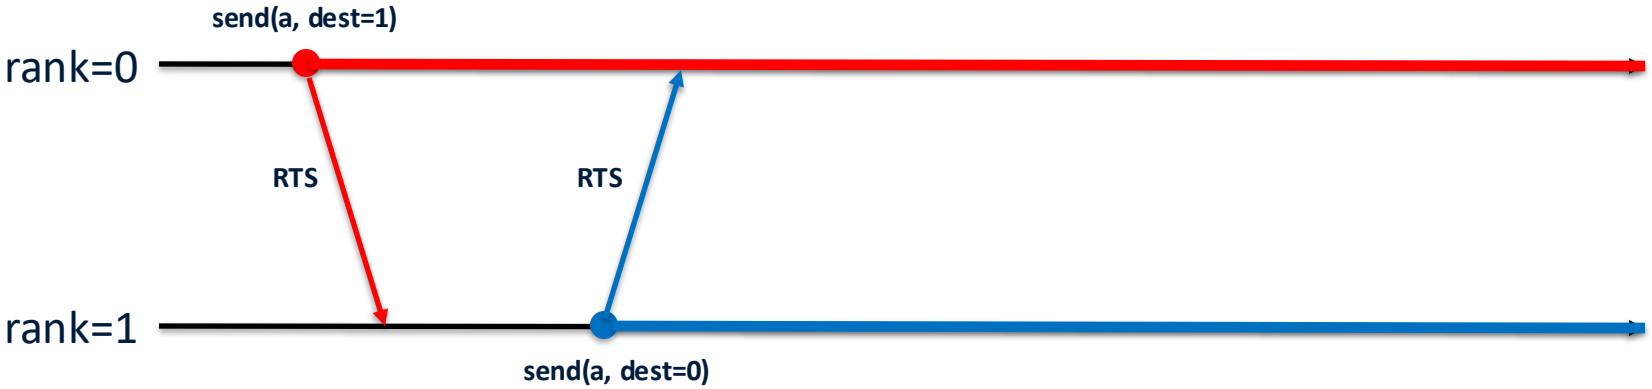

Matching Send / Recv

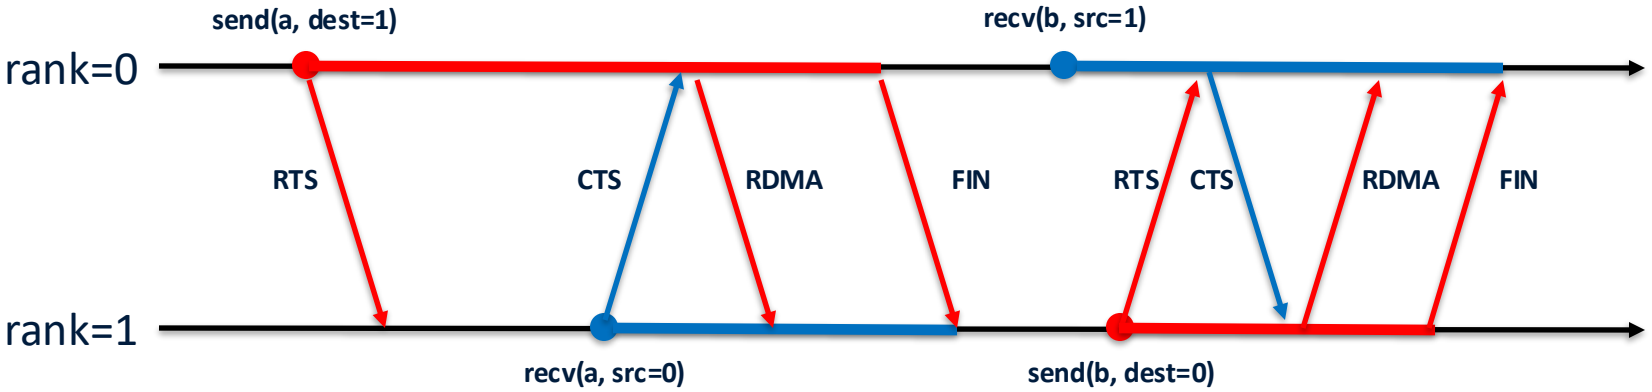

Buffering

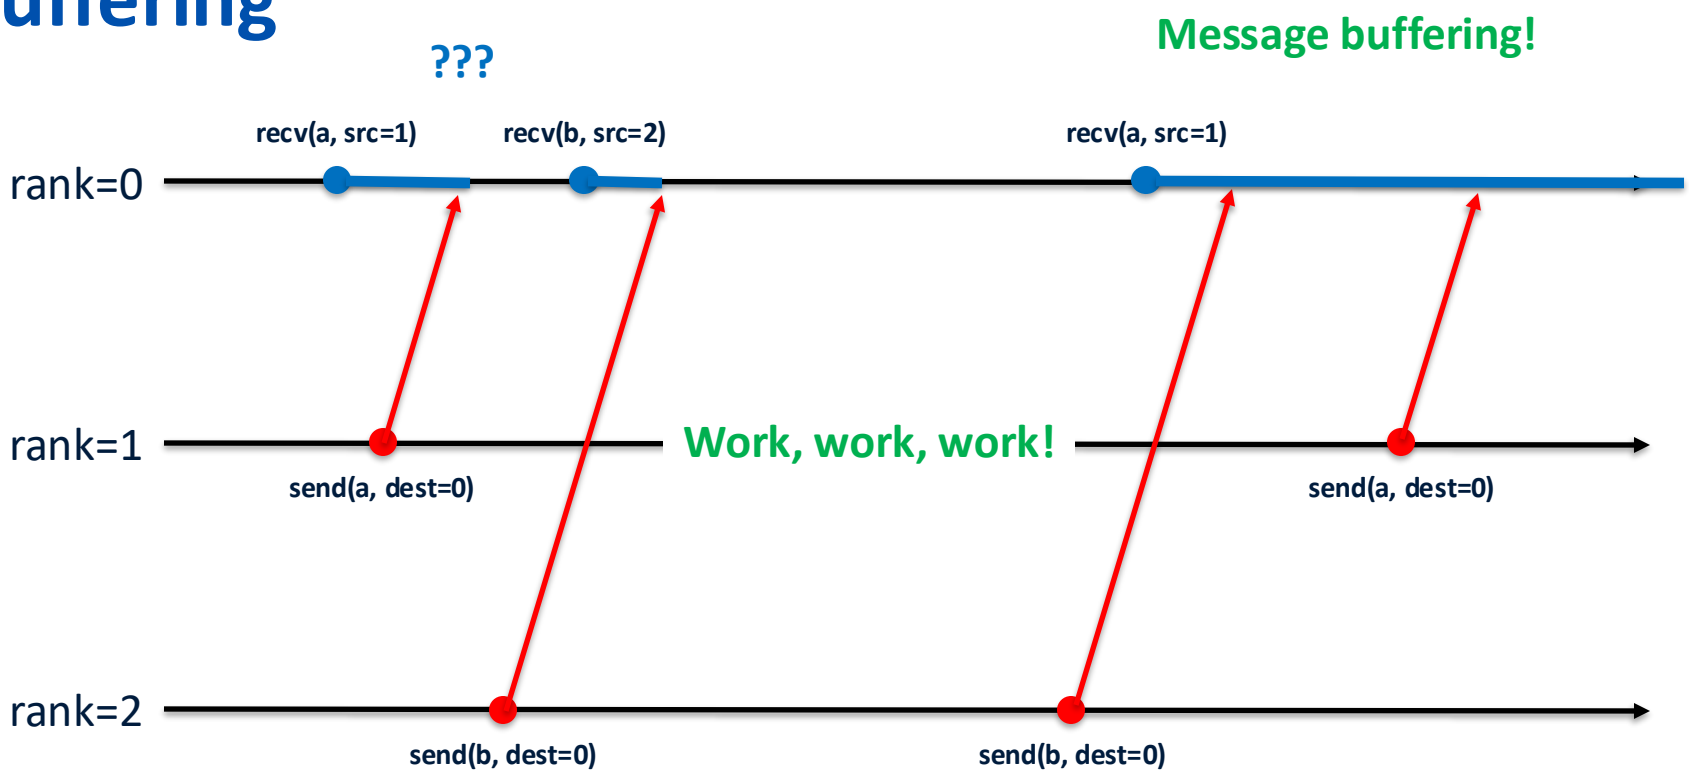
???
Message buffering!
recv(a, src=1)
recv(b, src=2)
recv(a, src=1)
rank=0
rank=1
rank=2
Work, work, work!
send(a, dest=0)
send(a, dest=0)
send(b, dest=0)
send(b, dest=0)

Tags

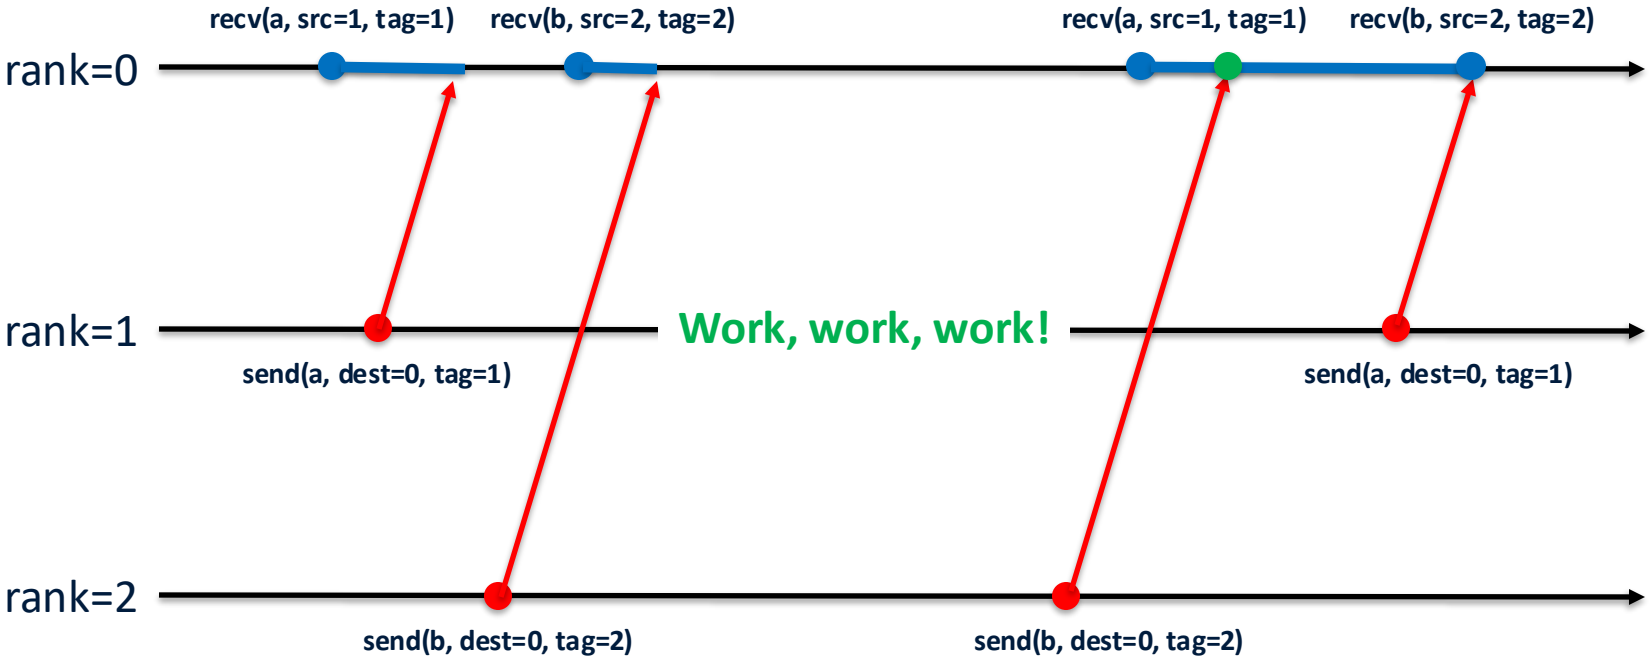
recv(a, src=1, tag=1)
recv(b, src=2, tag=2)
recv(a, src=1, tag=1)
recv(b, src=2, tag=2)
rank=0
rank=1
rank=2
Work, work, work!
send(a, dest=0, tag=1)
send(a, dest=0, tag=1)
send(b, dest=0, tag=2)
send(b, dest=0, tag=2)

Tags

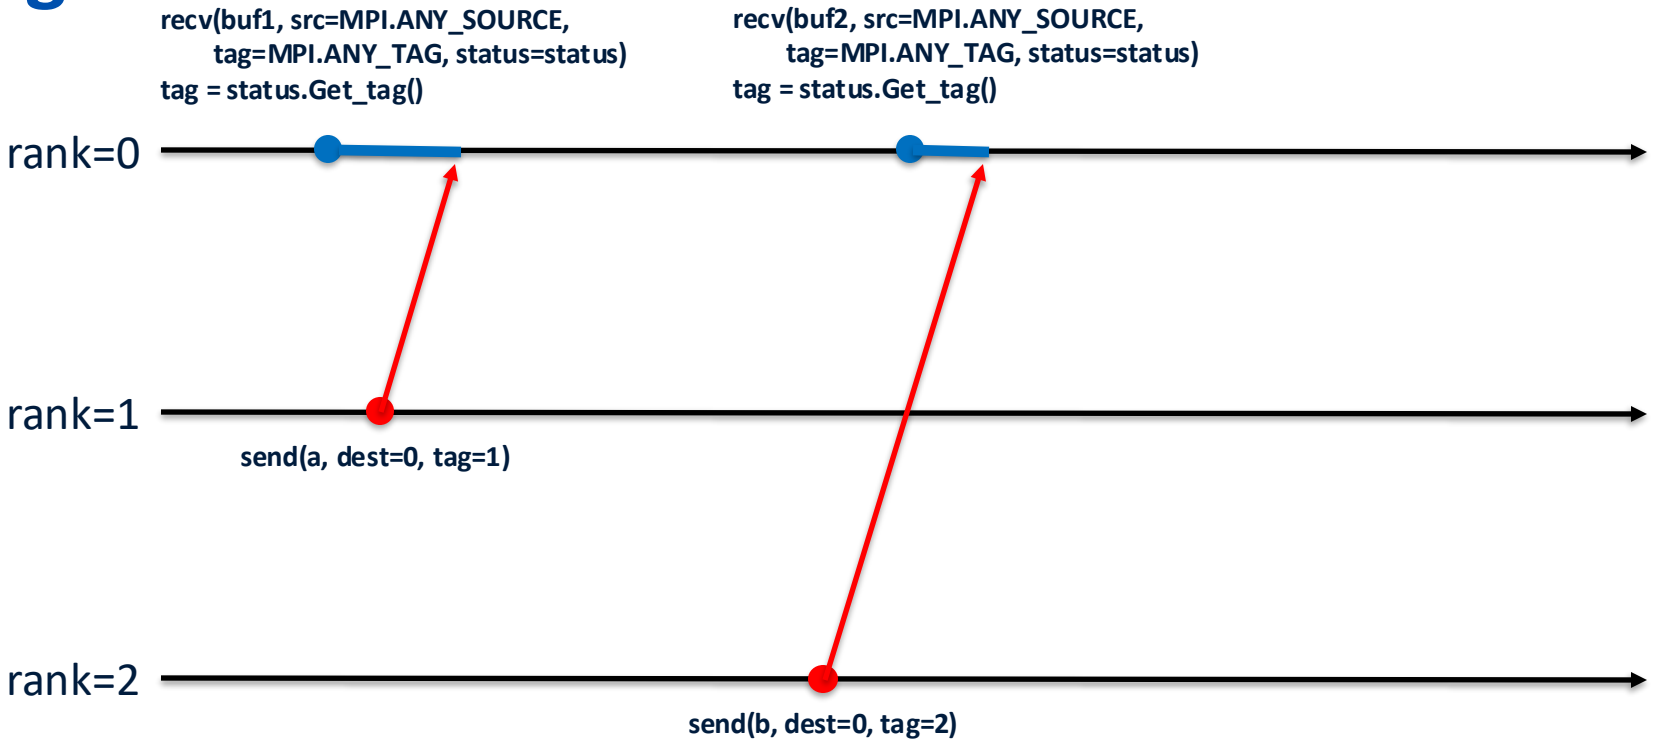
recv(buf1, src=MPI.ANY_SOURCE,
tag=MPI.ANY_TAG, status=status)
tag = status.Get_tag()
recv(buf2, src=MPI.ANY_SOURCE,
tag=MPI.ANY_TAG, status=status)
tag = status.Get_tag()
rank=0
rank=1
rank=2
send(a, dest=0, tag=1)
send(b, dest=0, tag=2)

Corners

global domain

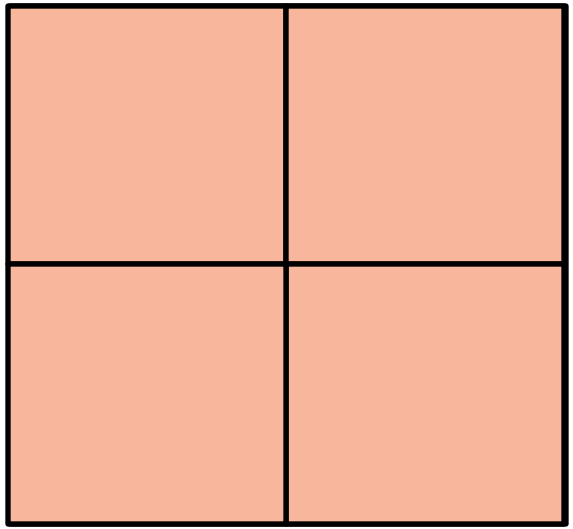

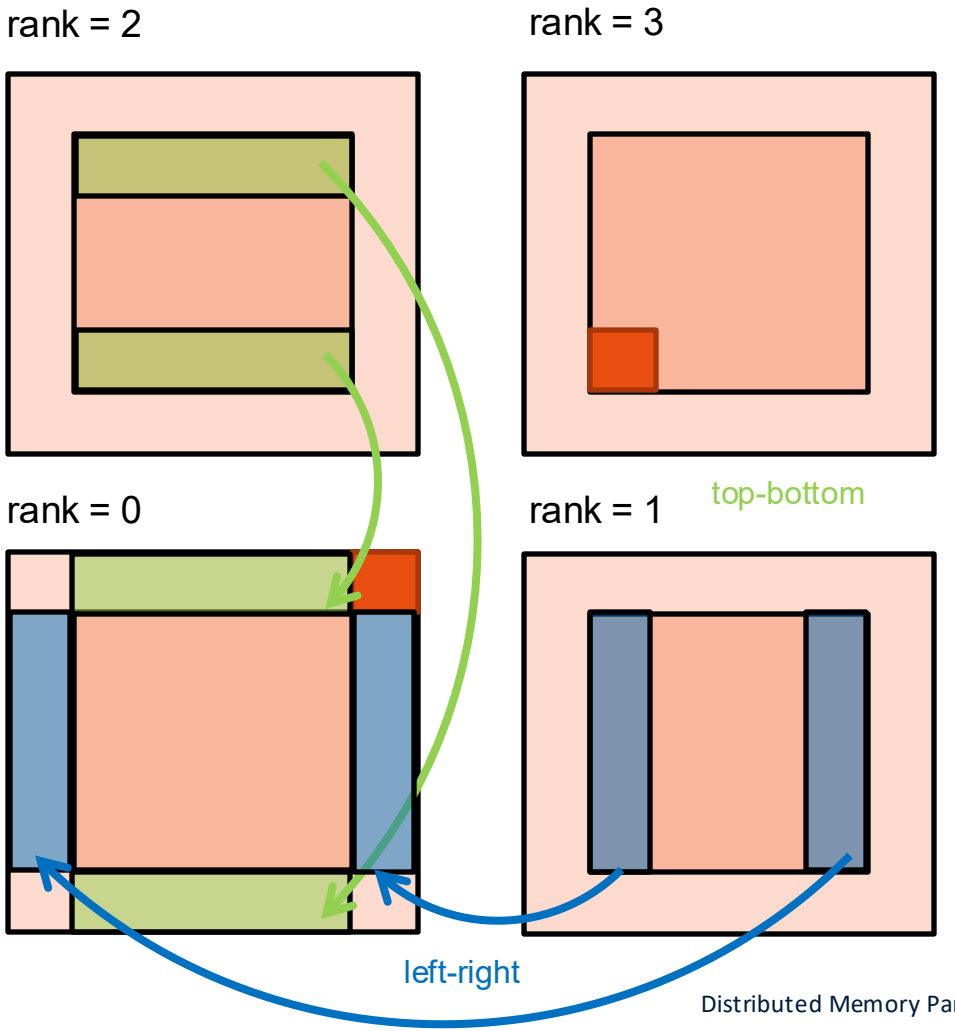

Corner Strategy 1

global domain

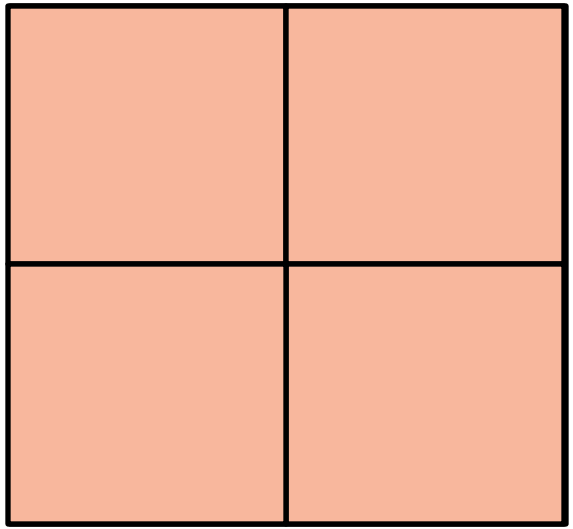

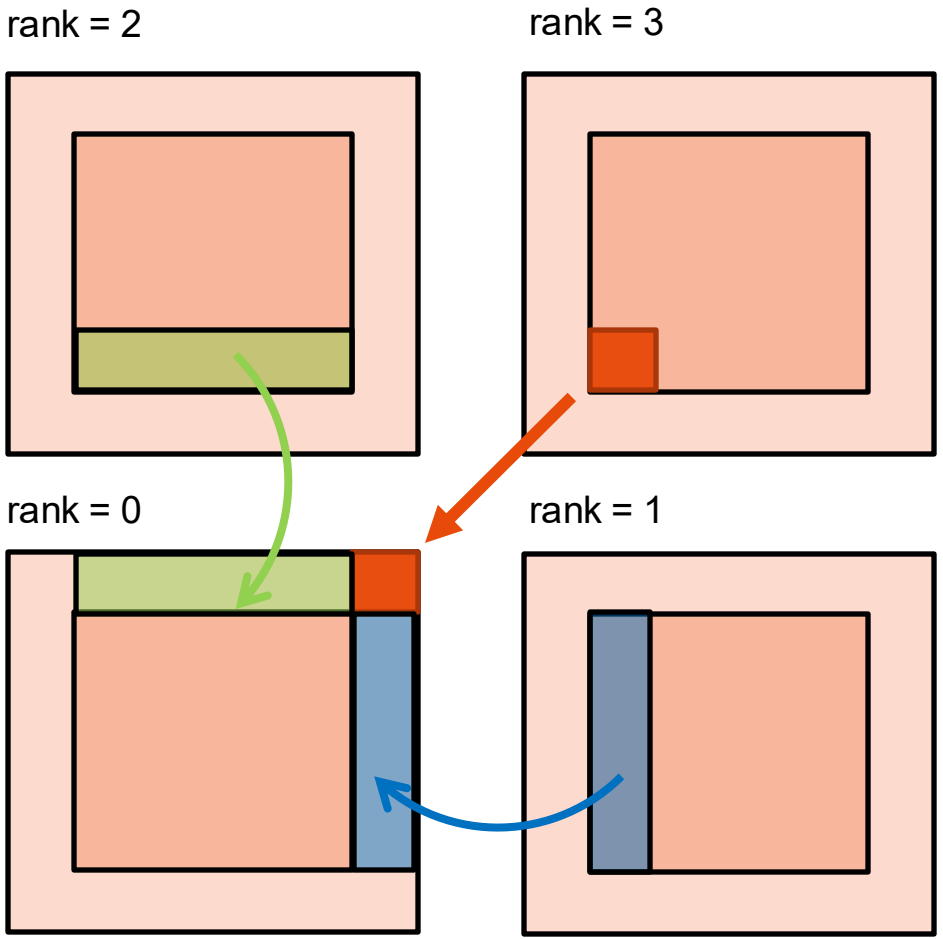

Corner Strategy 2

global domain

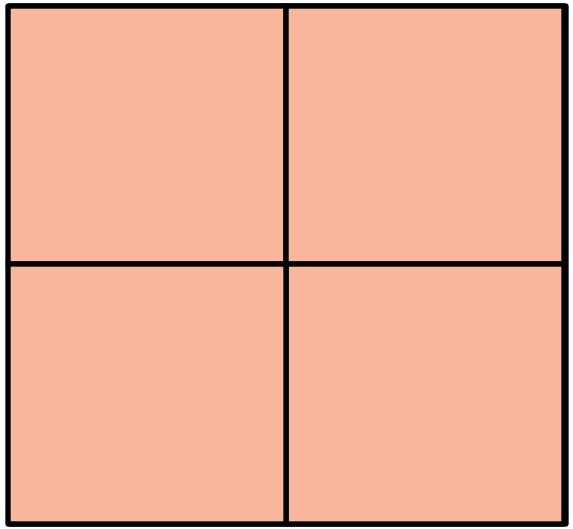

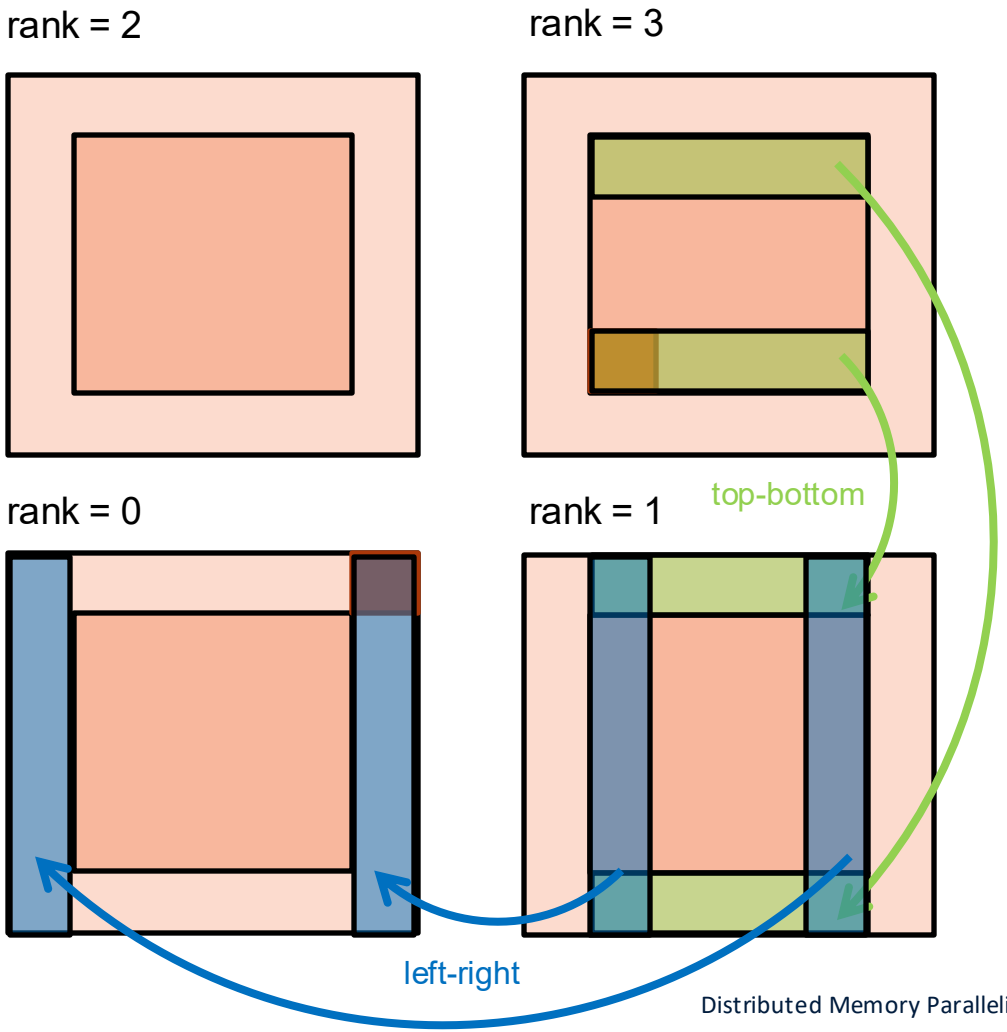

Error without sync

global domain